

Agentic Access to Brokerage Capabilities

Why a customer-facing MCP server is the wrong control surface for a retail broker-dealer — and the reference architecture that gives customers everything they are asking for without it.

RECOMMENDATION

Do not stand up a public, firm-hosted MCP server for the five brokerage domains. Build an Agentic Access Gateway over the existing capability APIs, and publish a signed reference MCP adapter that customers run inside their own agent runtime. Customers keep MCP. We keep the boundary.

PREPARED FOR	Chief Technology Officer; Chief Information Security Officer; Head of Brokerage Platform; Enterprise Architecture Council; Solution Architecture Guild
PREPARED BY	Technical Architect — Brokerage Trading APIs (Account Management, Customer Management, Trading, Investment Research, Money Movement)
SCOPE	External, customer-initiated agentic access to the five brokerage business domains by frontier models operated by third parties
DECISION SOUGHT	Endorsement of Option C and funding for Phase 0 (see §11–12)
CLASSIFICATION	Confidential — internal architecture review. Contains no customer data.

CONTENTS

What is in this document

Twelve sections and four appendices. Sections 1 and 12 are written to be read on their own; sections 5 through 9 carry the architecture; sections 3 and 10 carry the argument. Twenty-one figures do most of the work.

1	Executive summary	3
	The recommendation, the reasoning behind it, and what we are asking for	
2	What our customers are actually asking for	5
	Demand mapped across the five domains, separated into read and write	
3	Why a customer-facing MCP server is a high-severity risk	8
	The structural argument, the threat model, the risk register, and why the standard mitigations do not close the gap	
4	Reframing: the four decisions MCP silently bundles	19
	Transport, capability contract, discovery, and authorisation are separable — and should be separated	
5	Target-state reference architecture	22
	The Agentic Access Gateway, the capability contract layer, and the two supporting planes	
6	Security mechanisms in detail	26
	What to use to expose battle-tested APIs indirectly, and why each control earns its place	
7	What our developers should build	32
	Capabilities, skills, registry and adapters — and how the effort should be split	
8	Customer journeys	37
	Onboarding, the read path, and the gated write path, end to end	
9	Control walkthroughs	44
	Five sequence diagrams: grant, read, trade, injection, revocation	
10	Options considered	55
	Four options assessed against eight criteria	
11	Roadmap and maturity gates	58
	Four phases, each independently valuable, with an explicit gate to revisit MCP	
12	Risks of this recommendation, and what would change our mind	61
	The honest counter-case	
A-D	Appendices	63
	Control matrix notes · worked grant · glossary · references	

SECTION 1

Executive summary

Our retail customers are right to want their AI assistants connected to their brokerage accounts. They are wrong about the mechanism, and so are we if we simply give them what they asked for. The demand is for agentic access; MCP is one implementation of it, and it is the one implementation that puts the authorisation decision outside our supervisory perimeter.

1.1 The recommendation in one paragraph

Build an **Agentic Access Gateway** in front of the five existing business capability APIs: a single, governed HTTPS surface where identity, a bounded grant, a policy verdict, a spending budget and an immutable audit record all meet on every call. Publish the capability contract protocol-neutrally — OpenAPI 3.1 with strict JSON Schema, a signed capability registry, and a set of Agent Skills that teach any model our house method. Then ship a **signed reference MCP adapter that the customer runs inside their own runtime**. Customers who want MCP get MCP on day one. What they do not get is an MCP server operated by us on the public internet, because that artefact relocates the decision "should this order be placed?" into a probabilistic planner we neither build, version, test, nor supervise.

THE CORE FINDING

The risk in a customer-facing MCP server is not a vulnerability that a good implementation fixes. It is **structural**: the protocol's design point is that the model chooses the tool and the arguments. For a search assistant that is elegant. For an order-entry system it means the effective authorisation decision has been delegated to a component sitting on the wrong side of our trust boundary, steered by text we never see.

1.2 Six findings

#	FINDING	WHY IT MATTERS HERE
1	The authorisation decision moves outside our perimeter	In an MCP deployment the planner decides which capability runs, on which account, for what amount. We receive a tool name and arguments. Everything upstream of that — the prompt, the retrieved documents, the other connected servers, the model version — is invisible to us and unversioned by us.
2	Tool descriptions and tool responses are prompt-injection surface	The model reads our tool schemas as instructions in the same context as attacker-controlled web pages, filings and email bodies. OWASP ranks <i>Agent Goal Hijack</i> as ASI01 in its 2026 Top 10 for Agentic Applications precisely because no production system reliably separates instructions from data inside a context window.
3	MCP's strongest authorisation control does not apply to retail	Enterprise-Managed Authorization (ID-JAG) became a stable MCP extension in June 2026 and is genuinely good — but it works by putting a <i>corporate identity provider</i> in the middle. Retail brokerage customers do not have one. The specification itself frames the user-driven consent model as the appropriate one for consumer applications.
4	A scope cannot express a transaction	OAuth scopes are coarse by design. "Trading: write" is not the same permission as "sell up to \$40,000 of these three tickers, today, once". The gap between the two is where every material agentic loss will occur, and RFC 9396 already closes it — but only if we own the enforcement point.

5	We would still owe the regulator a supervisory system we cannot build	FINRA Rule 3110 is technology-neutral. FINRA's 2026 Annual Regulatory Oversight Report devotes its first standalone GenAI section to exactly this, naming <i>autonomy, scope and authority</i> , and <i>auditability and transparency</i> as the specific agentic risks. A record that begins at the last hop cannot reconstruct why an order was entered.
6	The performance risk is real and points the same way	Our APIs are excellent — for screens. A planner asked one portfolio question can emit dozens of fine-grained tool calls. The fix is coarse, intent-shaped capabilities, which is an argument for owning the contract rather than mechanically projecting existing endpoints into tools.

1.3 What the customer experience actually looks like

The uncomfortable truth for the "just ship MCP" position is that Option C is *better* customer experience, not worse. Under the recommendation:

- ▶ The customer installs a connector in one command, in whichever runtime they already use — Claude, ChatGPT, Gemini, Copilot, or something they wrote themselves.
- ▶ They authenticate on *our* domain with a passkey, and set a mandate in plain English: which accounts, which domains, what daily cap, what expiry. Read-only for ninety days is one click.
- ▶ Read and research work instantly and feel unguarded, because every guard already executed inside the gateway's latency budget.
- ▶ Trades and transfers require one biometric tap in our app, on a screen we render from a server-side proposal. Roughly twenty seconds, only on writes.
- ▶ They see an activity feed of everything the agent asked for, received and was refused — and can revoke in one tap.

None of that is available in an MCP deployment where the consent screen, the confirmation prompt and the tool inventory all live in software we do not control.

1.4 What we are asking for

Decision Endorse Option C (§10) as the target architecture, and record the decision not to expose a firm-hosted public MCP server ahead of the Phase 3 gate.	Funding Phase 0 only: the gateway, the FAPI 2.0 authorisation profile, the grant registry, intent receipts, and the first five read capabilities. Two quarters.	Mandate A standing platform team to own the gateway, registry, adapters and SDKs, and a security engineering commitment to the injection corpus in CI.
--	---	--

Phase 3 deliberately reopens the MCP question with evidence rather than opinion. If, by late 2027, the protocol has a retail-grade consent model, attested clients and a confirmation channel we can observe, hosting our own MCP surface becomes a twelve-week adapter project on top of everything we will have built — precisely because none of the controls will have been buried in the protocol layer.

SECTION 2

What our customers are actually asking for

Before arguing about protocols it is worth being precise about the demand. Customers are not asking for MCP. They are asking for their assistant to be able to see and act on their money. MCP is simply the word they have learned for it.

2.1 The demand, mapped

Figure 1 maps the jobs customers describe onto the five domains, split by whether they read state or change it. Two things stand out. First, the overwhelming majority of the requests — and almost all of the delight — are reads. Second, the writes cluster in exactly the two domains where a mistake is expensive and irreversible: Trading and Money Movement.

FIGURE 1

Account Management	Customer Management	Trading	Investment Research	Money Movement
Positions, balances, cost basis, tax lots, statements, documents	Profile, suitability, KYC, beneficiaries, preferences, entitlements	Quotes, order entry, order state, executions, allocations	Screeners, fundamentals, estimates, ratings, third-party research	ACH, wires, internal transfers, journals, ACATS, payees
AGENT JOBS-TO-BE-DONE R Show me my consolidated position across all accounts R What is my unrealised gain and wash-sale exposure? R Retrieve my 1099-B and reconcile it W Open / re-title / close an account	AGENT JOBS-TO-BE-DONE W Update my address and phone W Add a beneficiary to my IRA R What is my risk tolerance on file? W Refresh my KYC attestation	AGENT JOBS-TO-BE-DONE W Place a limit order for 100 AAPL @ 210 W Cancel / replace my working orders W Rebalance to my target allocation R Show me today's fills and average price	AGENT JOBS-TO-BE-DONE R Screen for dividend growers under 18x earnings R Summarise the latest analyst notes on NVDA R Compare two ETFs on fees and tracking error R Back-test this idea against my holdings	AGENT JOBS-TO-BE-DONE W Move \$25,000 from settlement fund to my bank W Add a new external payee W Schedule a recurring monthly contribution R Show me pending transfers and their status
SENSITIVITY Confidential	SENSITIVITY Restricted (PII)	SENSITIVITY Restricted (market)	SENSITIVITY Internal / licensed	SENSITIVITY Restricted (funds)

Figure 1 — Agentic demand across the five brokerage domains. Read requests (green) dominate by volume and carry most of the perceived value; write requests (red) concentrate in Trading and Money Movement, where loss is immediate and difficult to unwind. Any design that treats these two populations identically is either too slow for the first or too permissive for the second.

2.2 What this tells us about the design

The asymmetry in Figure 1 is the single most useful fact in this report, because it means we do not face a uniform problem. We face two problems with opposite optimisation targets.

The read path must be invisible

High volume, low individual risk, extreme latency sensitivity. If a customer asks "how exposed am I to semiconductors?" and waits four seconds, the feature has failed regardless of how well governed it is. Controls here must execute inside the request, not around it.

The write path must be deliberate

Low volume, high individual risk, latency-insensitive. Twenty seconds of friction on a \$40,000 order is not a defect; it is the product. The mistake is applying that friction to the read path, or omitting it from the write path because the protocol makes it awkward.

This is also why "just expose the existing APIs as MCP tools" is the wrong shape of answer. It applies one uniform control posture — a bearer token and a scope — to both populations at once, and it does so at the granularity our existing APIs happen to have, which was chosen for a screen rendering a page, not for a planner answering a question.

2.3 A note on who the customer is

It matters that this is a *retail* population. Three consequences follow, and each of them removes a control that would otherwise be available:

PROPERTY OF THE POPULATION	CONSEQUENCE FOR THE ARCHITECTURE
No corporate identity provider	The enterprise-managed authorisation model that MCP standardised in 2026 — where an IdP centrally provisions and revokes connector access — is unavailable. There is no admin to configure policy, and no offboarding event to trigger revocation.
No procurement or vendor review	Customers will connect whatever their assistant offers, alongside whatever else they have already connected. We cannot assume any diligence on the client side, and we must assume co-residency with untrusted servers.
Heterogeneous, fast-moving runtimes	The same customer may use three different assistants in a month, and each will update underneath them. Any control we place in the client is a control with an unknown version and an unknown lifetime.

Everything in Section 3 follows from these three facts far more than from any defect in the MCP specification itself. The specification is well written; it is simply solving a different problem than ours.

SECTION 3

Why a customer-facing MCP server is a high-severity risk

This section makes the case in five moves: the structural argument, the threat model, the risk register, an honest assessment of the mitigations, and the regulatory and performance exposure. The argument is not that MCP is badly designed. It is that its design point and our obligations are in direct tension.

3.1 The structural argument

Every security architecture eventually reduces to one question: *where is the decision made, and who can be held to it?* In a conventional brokerage channel the answer is clean. The customer expresses intent by manipulating a control we rendered; the browser or app transmits a request whose shape we defined; our server authorises it against a session we issued; our books and records retain the whole chain. Every link is ours.

MCP deliberately breaks that chain, and does so for good reasons. Its premise is that the *model* should choose which tool to call and with what arguments, because that is what makes an assistant useful across a hundred different services without bespoke integration. The consequence is that the component making the operative decision is a probabilistic planner running inside a third party's infrastructure, on a model version that can change without notice, primed by a context window we cannot inspect and did not compose.

THE REFRAME THAT CLARIFIES EVERYTHING

A tool call arriving at our MCP server is not a customer instruction. It is *a model's inference about what the customer probably wants*, rendered into a schema-valid request. Treating it as an instruction — which is precisely what a bearer token plus a scope does — is the category error at the heart of the risk.


This is not a hypothetical concern that better engineering resolves. A well-implemented MCP server, fully compliant with the 2025-11-25 authorisation specification, still receives exactly one artefact: a tool name and a set of arguments, accompanied by a token whose scope was granted at some earlier time for an unspecified purpose. There is no field in the protocol that carries "the customer said this", and no field that carries "the customer confirmed this on a screen you rendered".

3.2 The threat model

Figure 2 draws the trust boundaries explicitly. It is worth studying, because most discussions of MCP security implicitly assume the boundary sits where the TLS connection terminates. It does not. The security-relevant boundary sits much further left, at the edge of the context window.

FIGURE 2

Trust boundaries in a customer-facing MCP deployment



Every element left of the firm boundary sits outside our supervisory perimeter, yet fully determines what is asked of our APIs.

Figure 2 — Trust boundaries in a customer-facing MCP deployment. Only the rightmost zone is ours. The de facto authorisation decision is taken in the middle zone, using inputs drawn from the untrusted content plane at the bottom. Our tool descriptions, our tool responses, an attacker's web page and the customer's genuine request all occupy the same undifferentiated context.

The four properties that make this hard

- ▶ **We cannot attest the client.** MCP supports dynamic client registration, and consumer runtimes update continuously. We cannot know which software version composed the call, nor which model planned it.
- ▶ **We cannot see the context.** Our tool descriptions sit alongside descriptions from every other connected server. A malicious co-resident server's description can shape how the model uses *our* tools — the tool-shadowing problem OWASP documents in its MCP Security Cheat Sheet.
- ▶ **Our own responses become attack surface.** Every response we return re-enters the context as text. If a research summary we serve contains adversarial content — and third-party research and filings routinely contain arbitrary text — we have carried the injection ourselves.
- ▶ **The confirmation lives in the client.** Both OWASP and the MCP specification require human-in-the-loop approval for sensitive actions. Both place that requirement on the *client*. A confirmation we cannot observe is not a confirmation we can evidence.

3.3 The risk register

Figure 3 plots twelve risks specific to a customer-facing MCP deployment of these five domains. Placement reflects our own assessment over a twenty-four month horizon, assuming a competent implementation that follows the current specification and the OWASP guidance in full.

FIGURE 3




Figure 3 — Risk register for a firm-hosted, customer-facing MCP server. Seven risks land in the high band. Critically, the four highest — injection-driven orders, tool-description poisoning, consent bypass, and unreconstructable audit — are not addressed by any control the protocol currently offers to a retail deployment.

Two risks worth expanding

Risk 1 — indirect prompt injection drives an unintended order. This is ASI01 (*Agent Goal Hijack*) in the OWASP Top 10 for Agentic Applications, published December 2025. The attack requires no compromise of our systems: an attacker places instructions in content the model will read during ordinary research. The model, unable to reliably distinguish that text from the customer's request, plans a capability call. If the grant permits it, we execute it, correctly, with a valid token, and record it as a legitimate customer action. Every log we hold will say it was.

Risk 8 — supply chain. In September 2025 a package named `postmark-mcp` on npm shipped fifteen clean releases mirroring an official repository, then in version 1.0.16 added a silent blind-copy of every outgoing email to an attacker-controlled domain. It was widely installed before disclosure. The lesson is not that MCP is uniquely vulnerable to supply-chain attack — it is that agentic connectors are installed by end users, at scale, without review, and that a trusted component can become hostile between one release and the next. In a brokerage context, the equivalent artefact would sit between our customer and their money.

3.4 Why the standard mitigations do not close the gap

It would be dishonest to present MCP as unmitigated. The specification has improved substantially: the June 2025 revision classified servers as OAuth 2.0 resource servers and mandated Resource Indicators (RFC 8707); the November 2025 revision added OpenID Connect discovery, incremental scope consent via `WWW-Authenticate`, and Client ID Metadata Documents. The June 2026 Enterprise-Managed Authorization extension is a genuinely strong control. Figure 6 assesses each against our specific exposure.

FIGURE 6

MCP CONTROL	WHAT IT ACTUALLY DOES	RESIDUAL GAP FOR A RETAIL BROKERAGE	GAP CLOSED?
OAuth 2.1 + PKCE, MCP as resource server	Spec 2025-06-18 makes the server an OAuth 2.0 resource server and requires PKCE.	Closes credential theft in the redirect. Says nothing about whether the authenticated principal actually intended this call.	Partly
Resource Indicators (RFC 8707) + audience validation	Tokens must be issued for, and validated by, this specific server.	Closes token passthrough and audience confusion. The token still carries a scope, not a transaction.	Partly
Protected Resource Metadata (RFC 9728) discovery	Server advertises its authorisation server; removes ambiguity.	Good hygiene. Irrelevant to whether the order was the customer's idea.	Partly
Incremental scope consent via WWW-Authenticate	2025-11-25 adds step-up scope challenges rather than up-front omnibus consent.	Real improvement over `full-access`. Still coarse: a scope cannot express "\$25,000, to this payee, once".	Partly
Enterprise-Managed Authorization (ID-JAG / XAA)	Stable June 2026. An enterprise IdP centrally provisions and revokes MCP access.	The strongest control MCP has — and it does not apply. Retail customers have no corporate IdP. The spec itself scopes this to enterprise environments.	No
Human-in-the-loop confirmation in the client	OWASP and the MCP spec both require explicit approval for sensitive actions.	Enforced by the client, which we do not build. A confirmation we cannot observe is not a confirmation we can evidence.	No
Tool-definition pinning / hash verification	Detects rug-pulls by re-hashing definitions before execution.	Protects the client from us. Does nothing to protect us from a poisoned co-resident server in the same context.	No
Strict JSON Schema, additionalProperties: false	Constrains the shape of arguments the model may send.	Blocks malformed calls. A perfectly well-formed sell order for the wrong reason still passes.	No
Sandboxing and least-privilege server credentials	Limits the blast radius of a compromised server process.	Sound practice, adopted in the recommendation. Orthogonal to the intent-integrity problem.	Partly

Figure 6 — MCP's own controls, assessed against a retail brokerage threat model. Every control does something real. None of them addresses intent integrity — whether the request reflects what the customer actually wanted — and the strongest of them assumes an enterprise identity provider that retail customers do not have.

THE PATTERN IN FIGURE 6


Read down the right-hand column and a pattern emerges. Every mitigation that MCP offers protects the *channel*: is the token real, is the audience right, is the client who it claims to be, has the definition changed. Not one of them protects the *intent*: is this what the human wanted. That is not an oversight in the specification — it is a problem the protocol layer cannot solve, because the protocol layer does not contain the human.

3.5 Regulatory exposure

The regulatory position is unusually clear, largely because FINRA has said so directly. Its 2026 Annual Regulatory Oversight Report contains the first standalone GenAI section in the report's history and states plainly that FINRA's rules, being technology-neutral, "continue to apply when firms use GenAI or similar technologies". The report's *Emerging Trends in GenAI: Agents* subsection then lists the specific risks: agents acting autonomously without human validation; agents acting beyond the user's actual or intended scope and authority; multi-step reasoning that makes outcomes difficult to trace, complicating auditability; and agents operating on sensitive data.

Those are, almost word for word, the four properties of a customer-facing MCP deployment described in §3.2. Figure 7 maps each obligation onto the architectural capability it forces.

FIGURE 7



What each obligation actually demands

OBLIGATION	THE CONCRETE DEMAND ON THIS ARCHITECTURE
FINRA Rule 3110 — Supervision	A supervisory system reasonably designed to achieve compliance, tailored to the business. For an agentic channel this means written procedures covering agent onboarding, grant issuance, monitoring, escalation and revocation — and a system that can actually enforce them. FINRA Regulatory Notice 24-09 had already reminded members that these obligations attach to generative AI use.
FINRA Rule 4511 and SEA Rule 17a-4 — Books and records	Records with no otherwise specified retention period must be preserved for at least six years. For agentic orders the record that matters is not the order ticket alone but the chain that produced it: intent, grant, policy verdict, confirmation, execution. If we cannot produce that chain, we have an order we cannot explain.
SEC Regulation S-P, as amended May 2024	A written incident-response programme, with notification to affected individuals as soon as practicable and no later than thirty days. Compliance dates fell in December 2025 for larger entities and June 2026 for smaller ones. An agentic channel materially expands the population of systems from which customer information can be exfiltrated — including through entirely legitimate-looking tool arguments.
Reg BI and suitability	If our capability surface shapes what the customer is told — through what we return, how we summarise it, or what a skill instructs the model to emphasise — we should expect that output to be assessed as ours. This argues strongly for owning the skills layer rather than leaving the model to invent house method.
BSA/AML and Regulation S-ID	Agentic money movement must feed the same red-flag and transaction-monitoring programmes as every other channel. This is straightforward if the channel terminates at a gateway we operate, and close to impossible if the decision path begins in a third party's runtime.

A CAVEAT ON REGULATION SCI

An earlier proposal would have extended Regulation SCI to broker-dealers above certain thresholds; the SEC withdrew that proposal in June 2025. We should therefore not lean on Reg SCI as a driver unless a specific entity in the group is already an SCI entity. It is mentioned here only to note that we have deliberately excluded it from the argument.

3.6 The performance and operational case

The performance concern raised in the original brief is well founded, but the mechanism is usually misdiagnosed. The problem is not that MCP's transport is slow — JSON-RPC over streamable HTTP is fine. The problem is *call amplification*, and it is a direct consequence of exposing existing API granularity as tools.

1 → 40

A single question such as "am I drifting from my target allocation?" can decompose into dozens of tool calls if positions, prices, classifications and targets are separate tools

xN

Each call is a separate authenticated round trip with its own token validation, policy evaluation and audit write — costs our screen-oriented APIs never paid

Unbounded

Retry and re-planning loops mean the call count is not a function of the customer's request but of the model's uncertainty, which we cannot cap from outside

Three second-order effects follow, and each of them is an operational risk in its own right:

- **The load signature is indistinguishable from abuse.** Bursty, high-fan-out, low-think-time traffic from a small number of authenticated principals is exactly what our existing rate-limiting heuristics are tuned to suppress. Either we suppress legitimate agents or we blind the control.
- **Cost becomes unpredictable per customer.** Where a screen made a bounded number of calls, an agent makes a distribution. Without a per-grant budget the marginal cost of a curious customer is unbounded.
- **Latency compounds.** Forty sequential calls at 40 ms is 1.6 seconds of pure network time before the model has said anything. The customer experiences this as the feature being slow, and blames us rather than their assistant.

The remedy — coarse, intent-shaped capabilities designed for a planner rather than a page — is the same remedy the security argument produces. That convergence is not a coincidence: both problems come from letting a protocol dictate the granularity of a business contract.

SECTION 4

Reframing: the four decisions MCP silently bundles

Most of the difficulty in this debate comes from a single conflation. "Should we stand up an MCP server?" sounds like one decision. It is four, and they have different answers.

FIGURE 4




Figure 4 — Deconstructing the question. Transport is a commodity. The capability contract is the durable asset. Discovery and guidance are a product surface. Authorisation and consent are a regulated control. Bundling them into one artefact means the weakest of the four determines the risk of all of them.

4.1 Working through the four

Transport and framing is genuinely a commodity, and we should treat it as such. JSON-RPC 2.0 over streamable HTTP is a reasonable choice; so is plain HTTPS with JSON bodies. The decision that matters is not which one we pick but *who runs the client half*. If the customer's runtime runs the MCP client and our reference adapter, we get MCP compatibility without operating an MCP server on the public internet.

The capability contract — the names, schemas, semantics, error taxonomy and data contracts of the operations we expose — is the real asset, and it should outlive every protocol currently in fashion. It should be expressed once, in a protocol-neutral form, versioned semantically, tested adversarially, and projected mechanically into whatever the market wants this year.

Discovery and guidance is a product surface, not a plumbing concern. An agent needs to find a capability and then needs to know how to use it well — our definitions of "drift", our conventions for tax lots, our rules about when a screen is and is not suitable. That is what Agent Skills are for, and they are considerably richer than a tool list.

Authorisation and consent is a regulated control and must never leave our side of the boundary. This is the one of the four where the answer is categorical rather than pragmatic.

THE CONSEQUENCE

Once the four are separated, the apparent trade-off between "give customers what they want" and "keep control" mostly dissolves. Customers want the first three. The risk lives almost entirely in the fourth. We can grant all of the first three, generously and quickly, provided we never bundle the fourth in with them.

4.2 What we are actually deciding

Stated precisely, the decision before the council is *not* whether to support MCP. It is whether the boundary at which authorisation is enforced sits inside our estate or inside the customer's runtime. Everything else in this report is an elaboration of that single choice.

SECTION 5

Target-state reference architecture

One public surface. Two supporting planes. The existing domain APIs untouched. The design is deliberately unexciting — its merit is that every control has exactly one home, and that home is on our side of the boundary.

FIGURE 5




Figure 5 – The Agentic Access Gateway and its supporting planes. Customer-owned runtimes sit entirely outside the trust boundary. The gateway is the sole public surface and the sole point at which identity, grant, policy, budget and audit meet. The five business capability APIs are unchanged.

5.1 The gateway

The Agentic Access Gateway is a policy enforcement point, not a proxy. Five functions live in it, and they must live together because each one depends on the others' output:

FUNCTION	RESPONSIBILITY
Channel security	Terminates mutual TLS or validates DPoP proofs. Every token is sender-constrained, so a stolen token is inert without the corresponding key. This is a FAPI 2.0 requirement and it is not negotiable for a money-moving surface.
Token and audience validation	Validates that the token was issued <i>for this gateway</i> (RFC 8707 resource indicators), is unexpired, is bound to the presented key, and is not on the revocation deny-list. Never forwards a customer-facing credential downstream.
Policy Decision Point	Evaluates attribute-based policy per call: the grant, the account set, the product, the notional and velocity budget, the step-up state, the customer's own restrictions. Returns a permit with <i>obligations</i> — redaction rules, row caps, watermarks — never a bare yes.
Content firewall	Inspects inbound arguments and outbound payloads. Inbound: schema conformance, argument plausibility, known injection signatures. Outbound: data-loss patterns, licence-bound research content, and instruction-shaped text in anything we return to a model.
Budget governor	Enforces per-grant, per-agent and per-capability quotas on call rate, data volume and notional. This is simultaneously a cost control, a performance control and a blast-radius control.

5.2 The capability contract layer

Between the gateway and the domain APIs sits the layer that does the intellectual work: turning screen-shaped endpoints into planner-shaped capabilities. Four components:

- **Capability registry** — a signed catalogue of every exposed operation with its owner, risk tier, required grant type, semantic version, changelog and deprecation date. This is the artefact an examiner will ask to see, and the artefact from which every adapter is generated.
- **Intent endpoints** — coarse, business-shaped operations. `portfolio.snapshot` rather than `getPosition` + `getPrices` + `getLots`. `trading.propose` and `trading.place` as two distinct capabilities with two distinct grants.
- **Schema and guardrails** — strict JSON Schema with `additionalProperties: false`, pattern constraints on every string, bounded numeric ranges, and enumerated instrument identifiers. This does not stop a well-formed bad request, but it eliminates the entire class of malformed ones cheaply.
- **Response shaper** — field-level redaction, data minimisation by default, provenance metadata (as-of timestamps, source capability, receipt identifier) attached to every response so that a hallucinated figure is falsifiable rather than merely plausible.

5.3 The identity, consent and authorisation plane

Six components, covered in detail in Section 6: the customer identity provider with step-up; a FAPI 2.0 authorisation server; a grant registry storing RFC 9396 `authorization_details`; an agent client registry holding software statements and attestations; a delegation service performing RFC 8693 token exchange; and a revocation service with three levels of blast radius.

5.4 The assurance, evidence and supervision plane

This plane is what converts the architecture from "secure" to "supervisable", and it is the part most likely to be under-resourced. It comprises the intent receipt service, the WORM audit store, a decision log capturing plan and policy verdict alongside outcome, live feeds into existing surveillance and AML monitoring, per-agent behavioural analytics, and — importantly — the customer-facing transparency surface. The last of these is not a compliance artefact; it is the control customers actually exercise, and the one that generates the trust the whole channel depends on.

5.5 What deliberately does not change

THE EXISTING APIS STAY AS THEY ARE

The five business capability APIs are battle-tested and performant, and nothing in this design asks them to change. They continue to serve the web and mobile channels on exactly the same code paths. The gateway and contract layer are additive. This matters for delivery risk: Phase 0 touches no system of record and no order path.

SECTION 6

Security mechanisms in detail

The brief asked directly: what security mechanisms should be used to expose these APIs indirectly? This section answers it control by control. Nothing here is novel — every mechanism is a published standard with production implementations. The novelty is only in assembling them into a chain where each artefact proves something the previous one cannot.

6.1 The chain of custody

FIGURE 8

Chain of custody — seven artefacts, each proving something the previous one cannot




Figure 8 — Seven artefacts, seven distinct proofs. In an MCP baseline, artefacts 1 through 4 collapse into a single bearer token whose scope was granted weeks earlier for an unspecified purpose, and artefacts 5 through 7 do not exist at all.

6.2 Baseline: FAPI 2.0, not plain OAuth

The starting point should be the OpenID Foundation's Financial-grade API 2.0 Security Profile rather than bare OAuth 2.1. FAPI 2.0 is a ground-up redesign that is simpler than FAPI 1.0 while raising the security bar, and it has been subjected to formal security analysis. Two of its requirements matter most here:

- **Sender-constrained tokens are mandatory**, via either mutual TLS or DPoP (RFC 9449). This single requirement removes bearer-token theft from the threat model. Given that agent runtimes store credentials on consumer devices, in configuration files, and occasionally in plaintext, this is the highest-value control we can adopt.
- **Pushed Authorization Requests** (RFC 9126) move the authorisation request off the browser URL and into a back-channel call, which is what makes it safe to carry a rich, structured grant request rather than a scope string.

We should also adopt RFC 9700 (OAuth 2.0 Security Best Current Practice) as the baseline for the authorisation server, and RFC 9728 Protected Resource Metadata for discovery, so that a client can locate our authorisation server without ambiguity.

6.3 The central mechanism: Rich Authorization Requests

If one control in this report deserves the most attention from architects, it is RFC 9396. Scopes are a flat namespace of strings; they can express "trading:write" and cannot express "sell up to \$40,000 of these three tickers, today, at most twice". RFC 9396 replaces the scope string with a structured `authorization_details` array carried through authorisation, consent, token issuance and enforcement.

The practical effect is that the customer's consent screen and the enforcement point speak the same language. What the customer approved is literally what the gateway evaluates. Appendix B gives a worked example.

WHY THIS CLOSES THE GAP FIGURE 6 IDENTIFIES

An injected instruction can cause a model to *request* anything. It cannot expand a grant. If the grant contains no money-movement authorisation detail, no amount of context manipulation produces a transfer — the request fails at the gateway with an insufficient-grant error, and the only path to elevation runs back through a consent screen that we render, on our domain, to the authenticated human.

6.4 Agent identity and client attestation

We should not accept dynamic client registration on this surface. Every agent that reaches production must be a named client with:

- a signed software statement identifying the publisher and the artefact;
- `private_key_jwt` client authentication, never a shared secret;
- exact-match redirect URIs, validated by string comparison rather than pattern;
- a declared capability set, so that a research client cannot request a trading grant;
- a revocable identity, so that a single disclosure can disable an entire agent family across all customers at once.

This is the control that would have contained the `postmark-mcp` class of incident: a compromised adapter version is a client-registry entry we can revoke centrally in one action, rather than a package on a million laptops.

6.5 Step-up and the out-of-band confirmation channel

This is the load-bearing control of the entire design, and it deserves to be stated plainly: **for any capability that moves money or changes ownership, authority is granted on a channel the firm renders, not in the model's context.**

Mechanically this is decoupled authorisation — the pattern standardised as CIBA (Client-Initiated Backchannel Authentication) in the OpenID family, and familiar from every bank push-approval flow. The important design details are:

DESIGN DETAIL	WHY IT IS ESSENTIAL
The payload is server-rendered	The customer sees a proposal our Trading API produced and signed, not a summary the model wrote. A compromised context can request anything; it cannot author what the customer reads.
The proposal identifier is opaque and single-use	Confirmation is bound to one specific priced, sized, validated order. It cannot be replayed against a different one, and it expires.
The confirmation is the record	The signed confirmation assertion is what the intent receipt binds to. This is what makes the order reconstructable years later under Rule 4511 retention.
Money movement requires a second device	For transfers, confirmation must not be granted on the same device that initiated the request, and a new payee triggers a cooling-off window before it becomes usable.

6.6 Downstream isolation

The domain APIs must never see a customer-facing credential. The gateway performs an RFC 8693 token exchange, minting a short-lived internal token whose `act` claim names the agent and whose subject remains the customer. Three benefits follow: the blast radius of a leaked external token stops at the gateway; downstream logs correctly attribute the call to the agentic channel; and the domain teams need no knowledge of the external authorisation model at all.

6.7 The control baseline, by domain

Figure 9 states the baseline control set for each domain. It is deliberately conservative on writes and deliberately permissive on reads, for the reasons set out in §2.2.

FIGURE 9

CONTROL BASELINE BY DOMAIN

	Read	Write	Grant scope	Sender-constrained token	Step-up auth	Out-of-band confirm	Notional / velocity cap	Cooling-off window	Surveillance feed	Intent receipt
Account Management Positions, balances, tax lots, documents	✓	○	✓	✓	○	○	—	—	○	✓
Customer Management Profile, KYC, beneficiaries, entitlements	✓	○	✓	✓	✓	✓	—	✓	✓	✓
Trading Quotes, order entry, cancel/replace, executions	✓	○	✓	✓	✓	✓	✓	○	✓	✓
Investment Research Screeners, fundamentals, licensed research	✓	—	✓	✓	—	—	○	—	○	○
Money Movement ACH, wires, journals, payees, ACATS	✓	○	✓	✓	✓	✓	✓	✓	✓	✓

KEY
  Required for every call
  Conditional — risk-tiered
  Not applicable

Read access is universal because reading is where the customer value is. Write access is conditional everywhere, and never straight-through from a model.

Figure 9 — Control baseline by domain. Filled marks are required on every call; hollow marks are conditional on risk tier. Read access is universal because reading is where the value is; write access is conditional everywhere and never straight-through from a model.

Notes on specific cells

- **Investment Research has no write path.** It is also the domain with the most valuable data and the tightest licensing constraints, which is why licence-aware response shaping and a data-volume budget matter more there than step-up does.
- **Customer Management is treated as high-risk despite not moving money.** Changing an address, adding a beneficiary or refreshing a KYC attestation are classic account-takeover precursors. An agent that can silently change a mailing address is an agent that can prepare a fraud.
- **Money Movement carries every control.** It is the only domain where we recommend both a cooling-off window and a second-device confirmation as unconditional requirements.

6.8 Observability and evidence

The intent receipt is the artefact that makes the channel supervisable. It is written *before* the acting call completes, not asynchronously afterwards, and it binds:

- the authenticated customer and the account set in play;
- the agent client identity and version;
- the grant identifier and the specific `authorization_details` element relied upon;
- the capability, its version, and the exact arguments after schema validation;
- the policy verdict and every obligation applied;
- the confirmation assertion, where one was required;
- the downstream outcome, including any venue or payment-rail reference.

Receipts stream simultaneously to WORM storage for retention, to surveillance and AML monitoring for supervision, and to the customer activity feed for transparency. Denials are recorded with the same fidelity as approvals — a denial pattern across the agentic cohort is one of the earliest signals of an injection campaign, and it is invisible if we only log what succeeded.

SECTION 7

What our developers should build

The brief asked whether the team should focus on tools, a tool registry, tool help, or skills. The answer is all four, in that order of importance reversed — and with a fifth item, capabilities, that is more important than any of them.

FIGURE 18

Tools, skills, resources and prompts — what our developers should actually build

The most common failure mode is to answer “build tools” and stop. Tools are the smallest and least differentiated part of the answer.




Figure 18 — Four things to build, and roughly how the effort divides. Tools, in the narrow MCP sense, do not appear on this diagram: they are a projection of the capability contract, generated rather than authored.

7.1 Capabilities first

A capability is a versioned, intent-shaped operation over a business domain, with a strict schema, explicit error semantics and a stated data contract. The design rule is that a capability should correspond to something a customer would say, not to a row in a database.

INSTEAD OF EXPOSING...	EXPOSE...	BECAUSE
<code>getPositions</code> , <code>getPrices</code> , <code>getLots</code> , <code>getTargets</code>	<code>portfolio.snapshot</code>	One consistent, coherent, as-of-stamped view. Forty round trips become one. Consistency becomes a property of the response rather than a hope.
<code>placeOrder</code>	<code>trading.propose</code> + <code>trading.place</code>	Separates the part a model may do unaided from the part that requires human authority. The model can explore freely; it cannot act alone.
<code>searchSecurities</code> with 40 optional filters	<code>research.screen</code>	A bounded, enumerated filter set is auditable, cacheable and cheap. An open query surface is an injection and cost-amplification vector.
<code>createTransfer</code>	<code>transfer.initiate</code> against a payee allow-list	The destination becomes a reference to something the customer already approved, rather than a free-text field the model fills in.

7.2 Skills — the answer to "how should the model go about this?"

Agent Skills, released as an open standard in December 2025, are packaged instruction sets: a `SKILL.md` file with YAML front matter plus optional templates, reference material and executable scripts. They load by progressive disclosure — only the name and description are in context at discovery, and the full body loads when a task matches — which keeps the context cost per available skill very small. The format has been adopted well beyond a single vendor, across roughly forty clients including Copilot, Cursor, Codex and Gemini CLI.

For a brokerage, skills are where the firm's judgement lives, and they are strategically more interesting than tools because they are portable and they differentiate:

What a skill should contain

Our definitions (what "drift" means in our statements), our conventions (how we present tax lots), our sequencing (always price before proposing), our guardrails (never characterise a screen result as a recommendation), and worked examples the model can pattern-match against.

What a skill must never contain

Authorisation logic, limits, or anything that would be load-bearing if an attacker rewrote it. A skill is guidance, not a control. If a rule matters, it belongs in the Policy Decision Point, where it is enforced rather than suggested.

THE DISTINCTION THAT KEEPS TEAMS HONEST

MCP modularises *access to systems*. Skills modularise *method*. Teams that conflate the two end up encoding judgement in tool descriptions — which is exactly the text an attacker most wants to influence, and exactly the text that changes silently when a co-resident server is added.

7.3 The registry and machine-readable help

Discovery should be boring and standard. We publish a well-known capability manifest, an OpenAPI 3.1 document with strict JSON Schema, an `llms.txt` and `AGENTS.md` for agents that read documentation rather than schemas, and a changelog with semantic versions and dated deprecations. The registry entry for each capability carries its owner, risk tier, required grant type and control set — which means a capability that has not been bound to a control set cannot be published. The registry, not a wiki, is the enforcement artefact.

7.4 Adapters, generated not authored

Adapters are thin, disposable projections of the registry. We should generate and maintain four: a signed reference MCP adapter that the customer runs locally, an OpenAPI document suitable for direct function calling by any frontier model, typed SDKs for Python and TypeScript, and plain documented HTTPS for anyone who wants neither. When the protocol landscape shifts again — and it will — we regenerate adapters and change nothing else.

THIS IS HOW WE SAY YES TO THE MCP REQUEST

A customer who wants MCP installs our reference adapter into their own runtime with one command. From their point of view they now have an MCP connector to their brokerage: tools appear, autocomplete works, everything behaves as expected. The difference is architectural and invisible to them — the MCP client and server both live in their process, and the only thing crossing the public internet is an authenticated, sender-constrained, grant-bounded HTTPS call to a gateway we govern.

7.5 The lifecycle

FIGURE 19

Capability lifecycle — how a business capability becomes an agent-reachable, governed surface




Figure 19 — How a business capability becomes an agent-reachable, governed surface. Step 4 is the one most firms skip. An injection corpus that grows with every incident, run against every capability in CI, is the only defence in this architecture that improves over time.

SECTION 8

Customer journeys


The test of this architecture is not whether it is defensible in an audit. It is whether a customer who wanted "MCP for my brokerage" feels well served. These three journeys are the evidence.

8.1 Connecting an agent for the first time

FIGURE 10

Customer journey 1 — connecting an agent for the first time

The target is under seven minutes from “can it?” to a working, bounded, supervised connection — without a single sensitive byte moving before the mandate exists.



Design rule: the customer chooses the runtime; we choose the boundary. Every step above works identically whether they are running Claude, ChatGPT, Gemini or something they wrote themselves — because none of the controls live in the runtime.

Figure 10 — Day-zero onboarding. Target is under seven minutes from “can it?” to a working, bounded, supervised connection. Note that the only moment of real friction — setting the mandate — is also the moment that produces the artefact everything else depends on.

Two design choices in this journey are worth defending explicitly.

Authentication happens on our domain, not in the assistant. This costs a redirect. It buys phishing-resistant authentication on a surface we own, a consent screen we render, and a screenshot we can hand an examiner. Customers do not experience a passkey prompt on their broker's domain as friction; they experience it as reassurance, which is why the confidence curve rises at that stage rather than falling.


The mandate is expressed in English, not scopes. "Research and read-only, ninety days" is one click. "Trading up to \$10,000 a day, thirty days" is a deliberate second choice. RFC 9396 makes this possible without inventing a bespoke permission model, because the structured grant the customer approves is the same object the gateway later evaluates.

8.2 The read path — where the value is

FIGURE 11

Customer journey 2 — portfolio monitoring and research (read path)

The high-volume, high-value, low-risk journey. It should feel instant and unguarded, because every guard has already been applied at the boundary rather than in the customer's way.



This journey is where the customer's enthusiasm actually comes from. It is also the journey that a naive MCP deployment degrades most — dozens of small tool calls per question, each a separate authenticated round trip against APIs that were tuned for a screen, not a planner.

Figure 11 — Portfolio monitoring and research. On our working assumption this is the large majority of agentic traffic by volume. It should feel unguarded, which is only possible because every guard executes inside the request rather than around it.

This journey is where the customer's enthusiasm comes from, and it is also the journey a naive MCP deployment degrades most — dozens of small tool calls per question, each a separate authenticated round trip against APIs tuned for a screen. The single design decision that saves it is coarse, intent-shaped capabilities, which is the same decision the security argument produced in §3.6.

Note also the fourth stage. The analysis itself — reasoning about concentration, drift and tax lots — is the part frontier models are genuinely good at, and the part we should not try to own. Our contribution there is the skill that supplies method and house conventions, so the analysis is consistent with our own definitions rather than invented.

8.3 The gated write path

FIGURE 12

Customer journey 3 — an agent-initiated trade (gated write path)

The journey the firm is right to be nervous about. The design answer is not to refuse it, but to move the moment of authority out of the model's context and onto a channel we own.




Figure 12 — An agent-initiated trade. The dip in the confidence curve at stage three is intentional. Friction placed exactly where the risk is concentrated, and nowhere else, reads to a customer as competence rather than obstruction.

The journey the firm is right to be nervous about. The design answer is not to refuse it — refusing simply moves the demand somewhere we cannot see — but to move the moment of authority out of the model's context and onto a channel we own, render and record. Money movement follows the same six stages with three additions: a payee allow-list only the customer can extend, a mandatory cooling-off window on first use of a new destination, and confirmation from a different device than the one that initiated the request.

SECTION 9

Control walkthroughs

Five sequences. The first creates authority; the second and third spend it; the fourth shows what happens when someone tries to steal it; the fifth destroys it.

9.1 Registering an agent and issuing a bounded grant

Thirteen steps, of which the customer experiences three: a login, a consent screen, and a return to their assistant. Everything else is machinery. The flow is a conventional OAuth authorisation code exchange with four deliberate additions — a pushed authorization request so the grant can be structured rather than a scope string, a client-registry check before the consent screen renders, a step-up challenge whenever the request contains any write, and a persisted grant that outlives the token.

The detail most likely to be argued about in implementation is step 2. Refusing dynamic client registration means an agent publisher must be onboarded before their customers can connect, which is friction we are choosing to accept. The alternative — allowing any client to self-register against a money-moving surface — removes our ability to revoke an agent family in one action, and that is the control that contains a supply-chain incident.

FIGURE 13

Sequence 1 – registering an agent and issuing a bounded grant

The only moment at which authority is created. Everything later in the system is a consequence of what this flow wrote down.




Figure 13 – Grant issuance. The only moment at which authority is created. Everything later in the system is a consequence of what this flow wrote down – which is why the durable artefact is the grant, not the token.

9.2 A read call, end to end

Nine steps executing inside a target of twenty-five milliseconds of added latency at the ninety-fifth percentile. Two details are worth drawing out. First, the policy verdict at step 4 returns *obligations* rather than a bare permission — redaction rules and row caps travel with the permit and are applied by the gateway, so a domain API never has to know which channel it is serving. Second, the receipt at step 8 is written before the response is released, not asynchronously afterwards; a response that reaches the customer without a durable record of why is a books-and-records exception, and it should be impossible by construction rather than unlikely in practice.

FIGURE 14

Sequence 2 — a read call, end to end

The common case: roughly 95% of agentic traffic. It must be fast, cheap and fully evidenced at the same time.




Figure 14 — The common case. Nine steps, of which the customer perceives one. The whole control set executes inside the gateway's latency budget, which is what allows the read journey to feel unguarded while being the most heavily instrumented channel we operate.

9.3 A trade, with out-of-band confirmation

The split between `trading.propose` and `trading.place` is what makes this sequence work. Propose is safe to expose generously: it prices, sizes and validates an order through exactly the same code path the web channel uses, and returns a signed server-side proposal with an opaque identifier. The model may call it as often as it likes, reason about the result, and revise. It cannot place anything.

Place accepts only a proposal identifier and a confirmation assertion. Because the identifier refers to a proposal our own systems constructed, the customer's confirmation screen is rendered from our data rather than from anything the model produced — which is the property that makes a compromised context insufficient rather than merely inconvenient.

FIGURE 15

Sequence 3 — an agent-initiated trade with out-of-band confirmation

The central control of the whole design: the moment of authority is moved out of the model's context and onto a channel the firm owns, renders and records.




Figure 15 — The load-bearing control. Steps 6 through 9 move the moment of authority onto a channel the firm renders. A compromised context can request anything; it cannot author what the customer is shown, and it cannot produce the assertion.

9.4 An indirect prompt-injection attempt

The scenario is deliberately mundane: a customer asks their assistant to research a holding; the assistant fetches a page; the page contains text instructing it to move money. No system of ours is compromised, no credential is stolen, and the resulting request is perfectly well formed. This is the attack class OWASP ranks first among agentic risks, and it is the one for which no model-side defence can currently be relied upon.

Figure 16 shows four independent controls, any one of which stops the transfer. They are deliberately of different kinds — a grant constraint, a customer-managed allow-list, a human confirmation and a numeric cap — so that a single bypass technique cannot defeat more than one of them.

FIGURE 16

Sequence 4 — an indirect prompt-injection attempt, and the four places it fails

The attack the industry has no reliable model-side defence for. The architecture's answer is not to detect it, but to make a successful injection insufficient.




Figure 16 — Four independent failures. None of the four blocks prevents the injection. They prevent it from mattering — which is the only defence available, since no production system reliably separates instructions from data inside a context window.

READ THIS FIGURE HONESTLY

The correct summary of Figure 16 is not "we are protected against prompt injection". It is "we have arranged matters so that a successful prompt injection is insufficient". That distinction is the difference between a defence that degrades gracefully and one that fails completely the first time a model behaves unexpectedly. Any proposal that relies on the model resisting the instruction is relying on a control that does not exist.

9.5 Revocation and containment

Revocation is the control an examiner will ask about first and the one a customer will actually use, so it deserves to be a first-class service rather than an administrative procedure. Three properties matter: it must be reachable by the customer without contacting us, it must operate at three levels of blast radius, and its effect must be measured continuously rather than asserted annually.

FIGURE 17

Sequence 5 — revocation and containment

The control an examiner will ask about first, and the one a customer will actually use. It must work in seconds, at three levels of blast radius.




Figure 17 — Three levels of blast radius. The control an examiner will ask about first, and the one a customer will actually use. Each level needs a named owner, a documented recovery time objective, and a quarterly test.

Two details in this sequence are easy to under-build and expensive to retrofit. The first is that the deny-list must propagate to every gateway node, in every region, within seconds — a revocation that takes as long as a token lifetime is not a revocation. The second is that the response to a revoked token must be a machine-readable revocation reason, so that a well-behaved agent stops rather than entering a retry loop against a surface that will never accept it again.

SECTION 10

Options considered

Four options, eight criteria. The assessment is our own, and the criteria are drawn from the six capabilities in Figure 7 plus the operational concerns in §3.6.

FIGURE 20

Options considered



Scale: 1 bar = poor, 4 bars = strong. Assessed against the six capabilities in Figure 7 and the risk register in Figure 3.

Option C is not "MCP or not MCP". It is "where does the control surface live". Customers who want MCP still get MCP — they run our signed reference adapter inside their own runtime, and it speaks to a boundary we govern.

Figure 20 — Options assessment. Option B — a firm-hosted public MCP server — scores highest on reach and lowest on every control criterion. Option C achieves near-identical reach because the customer still gets an MCP connector; the difference is which side of the boundary it runs on.

10.1 Why Option B fails despite being the obvious answer

Option B is the option customers asked for, and it should be taken seriously rather than dismissed. It has real merits: the shortest time to market, the lowest integration burden on customers, and immediate compatibility with connector directories that drive discovery. If our five domains were read-only research, Option B would probably be the right answer.

It fails because of two properties in combination. First, the authorisation decision leaves our perimeter, and no control in the specification brings it back for a retail population (Figure 6). Second, the two highest-value domains for customers — Trading and Money Movement — are precisely the two where that failure is unrecoverable. A read served on a stale assumption is an inconvenience; an order placed on one is a loss, a complaint, an arbitration and a supervisory finding.

10.2 Why Option D is insufficient alone

A firm-hosted assistant inside our own app is a good product and we should build one regardless. It is not an answer to this brief. Customers are not asking for another chat window; they are asking to use the assistant they already use, with the context they already have, alongside their calendar and their email and their documents. Option D declines that request, and the demand will go to whoever accepts it.

10.3 Why Option C is not a compromise

THE POINT OF THE RECOMMENDATION

Option C is not "MCP, but less". It is **the same protocol, the same customer experience, and a different location for the control surface**. Customers who want MCP get MCP — they run our signed reference adapter inside their own runtime. What changes is that the authenticated, grant-bounded, sender-constrained call crossing the public internet arrives at a gateway we govern, rather than a protocol endpoint that receives a tool name and hopes.

SECTION 11

Roadmap and maturity gates

Four phases. Each is independently valuable and independently defensible; nothing in a later phase is a prerequisite for shipping an earlier one. Phase 3 deliberately reopens the question this report closes.

FIGURE 21

Delivery roadmap and maturity gates

Each phase is independently valuable and independently defensible. Nothing in a later phase is a prerequisite for shipping an earlier one.




Figure 21 — Delivery roadmap. Phase 0 ships read access across all five domains with the full control set already in place. Writes wait until the confirmation channel exists — not because writes are harder to build, but because they are indefensible without it.

11.1 Sequencing rationale

The sequencing is driven by one principle: *never ship a capability whose control set is incomplete, on the promise of completing it later*. Phase 0 therefore looks over-engineered for read-only access — a full FAPI 2.0 profile, sender-constrained tokens, a grant registry and intent receipts, all to serve portfolio snapshots. That is deliberate. Building the control set once, against low-risk traffic, means Phase 2 is a matter of adding capabilities to an architecture that already enforces them, rather than retrofitting enforcement into a live channel.

11.2 What we would measure

MEASURE	TARGET	WHY THIS ONE
Time from install to first successful read	Under 7 minutes, p50	The whole customer-experience argument for Option C rests on this number. If it drifts past fifteen minutes we have rebuilt the friction we were trying to avoid.
Gateway added latency	Under 25 ms, p95	Policy evaluation, content inspection and receipt writing must fit inside the request. Above this the read journey degrades noticeably.
Calls per customer question	Median under 3	The single best proxy for whether our capabilities are intent-shaped or merely CRUD with new names.
Revocation propagation	Under 2 seconds, p99	Continuously measured, not tested annually. A revocation slower than a token lifetime is theatre.
Denial rate by reason	Trended weekly	A rising insufficient-grant rate across the cohort is the earliest available signal of an injection campaign.
Receipt completeness	100%	Any acting call without a complete receipt is a books-and-records exception. This should alert, not report.

11.3 The Phase 3 gate

We should commit now to revisiting the MCP question in late 2027 against evidence rather than sentiment. Three specific conditions would change the recommendation:

- a consent model in the protocol expressive enough to carry a transaction rather than a scope, for a population without a corporate identity provider;
- client attestation strong enough that we can know which software and which model version composed a call;
- a confirmation channel whose invocation and outcome the resource server can observe and evidence, rather than trust.

If those arrive, hosting our own MCP surface becomes roughly a twelve-week adapter project on top of everything in Phases 0 through 2 — because none of the controls will have been built into the protocol layer in the first place. That optionality is itself a reason to prefer Option C.

SECTION 12

Risks of this recommendation, and what would change our mind

A report that only argues one way is a sales document. These are the strongest objections to Option C, stated as their proponents would state them, with our responses.

12.1 The objections

OBJECTION	RESPONSE
"You are betting against the industry standard. MCP has won; you will be building a bespoke thing nobody wants to integrate with."	Partly fair, and the reason we ship a reference MCP adapter rather than refusing the protocol. But the bet is narrower than it sounds: we are betting that <i>where authorisation is enforced</i> will matter more than which wire format carries the call. Every adapter we generate is disposable; the capability contract and the control set are not.
"A customer-run adapter is worse experience than a hosted connector. Discovery in connector directories drives adoption, and we lose that."	Real cost, and we should not pretend otherwise. Mitigations: a one-command install, presence in each runtime's directory pointing at our adapter, and a Phase 1 target of under seven minutes end to end. If measurement shows adoption materially lagging peers, that is legitimate evidence for the Phase 3 gate.
"You are building an authorisation server, a policy engine, a confirmation service and a registry. That is a platform, and platforms overrun."	Fair, and the reason Phase 0 is scoped to read-only. Every component listed has mature commercial and open-source implementations; none of it is research. The genuinely new work is the capability contract and the injection corpus, and both are needed under any option including B.
"Out-of-band confirmation will annoy customers into abandoning the feature."	Testable, and we should test it. Note that it applies only to writes, that customers already accept it for card payments and bank transfers, and that the alternative — an unconfirmed agentic order — is the failure mode most likely to end the programme entirely.
"This over-indexes on prompt injection. Model providers are improving instruction hierarchies and this will look paranoid in two years."	Possibly. But the control we would remove on that basis is the confirmation channel, and the cost of being wrong is asymmetric: an unnecessary tap versus an unauthorised transfer. We should revisit the strength of the claim at the Phase 3 gate, not the existence of the control.

12.2 Residual risks we are accepting

- **We cannot prevent bad advice.** Nothing here stops a model reasoning poorly about a customer's portfolio. Skills raise the floor; they do not guarantee quality. The mitigation is honest positioning and a clear boundary between our data and the model's interpretation of it.
- **Read-path exfiltration remains possible.** An injected instruction can cause an agent to read data it is entitled to read and then transmit it elsewhere in the customer's own runtime. Budgets, volume caps and anomaly detection bound this; they do not eliminate it. This risk exists identically under Option B.
- **We are adding a critical dependency.** The gateway becomes a tier-one system on the path to trading for a growing population. It needs the availability engineering, capacity planning and incident response that implies from Phase 0, not from Phase 2.
- **Adapter drift.** Customers will run old adapter versions. We need a supported-version policy, deprecation warnings surfaced in responses, and a hard floor below which the gateway refuses to serve.

CLOSING POSITION

Our customers are asking for something reasonable and we should give it to them quickly. The disagreement is not about whether to open the brokerage to agents — it is about whether the sentence "the customer authorised this trade" is something we can prove or merely something we assume. Option C lets us prove it, keeps every protocol option open, and asks the customer for exactly one extra tap on the small fraction of interactions where their money actually moves.

Prepared by the Brokerage Trading API architecture team. Comments and dissent are welcome and will be attached to the decision record.

APPENDIX A

Control matrix notes

Supplementary detail for Figure 9, for the architects who will implement it.

A.1 Risk tiering

TIER	DEFINITION	BASELINE CONTROLS BEYOND THE UNIVERSAL SET
T0 — Public	No customer data. Reference and metadata only.	Rate limiting only. No grant required.
T1 — Read, non-sensitive	Customer data with no direct identifiers.	Grant scope, sender-constrained token, volume budget, receipt.
T2 — Read, sensitive	PII, tax identifiers, full account detail, licensed research.	Adds field-level redaction obligations, licence-aware shaping, tighter volume caps, elevated anomaly sensitivity.
T3 — Write, reversible	Preferences, alerts, non-financial profile fields.	Adds step-up authentication and a notification to the customer.
T4 — Write, financial	Orders, cancels, transfers, beneficiary and address changes.	Adds out-of-band confirmation, notional and velocity caps, cooling-off where a new counterparty is introduced, surveillance feed, mandatory receipt before completion.

A.2 Obligations the Policy Decision Point may attach

- `redact` — named fields removed or masked before the response leaves the gateway.
- `row_cap` — maximum records returned, with an explicit truncation marker so the model does not silently reason over a partial set.
- `as_of_max_age` — refuse to serve data staler than the stated bound rather than serving it unmarked.
- `licence_class` — third-party research content tagged so downstream shaping can enforce vendor terms.
- `watermark` — provenance and receipt identifier embedded in the response for later attribution.
- `require_step_up` — returned as a challenge rather than a denial, so the agent can prompt the customer to elevate.

A.3 Failure semantics

Every control fails closed, and every denial carries a machine-readable reason so that a well-behaved agent can respond sensibly rather than retrying. The reason taxonomy should distinguish, at minimum: `insufficient_grant`, `grant_expired`, `budget_exhausted`, `confirmation_required`, `confirmation_declined`, `cooling_off`, `payee_not_allowed`, `step_up_required`, `capability_deprecated`, and `revoked`. These strings will end up in customer-facing text written by models we do not control, so they should be chosen to be safe when quoted verbatim.

APPENDIX B

A worked grant

What "bounded authority" looks like on the wire. This is the object the customer approves on our consent screen and the object the gateway evaluates on every call — the same object, which is the entire point of RFC 9396.

```
// Pushed authorization request (RFC 9126), authorization_details per RFC 9396
{
  "client_id": "agent.acme-assistant.v3",
  "response_type": "code",
  "code_challenge_method": "S256",
  "authorization_details": [
    {
      "type": "brokerage_read",
      "domains": ["account_management", "investment_research"],
      "accounts": ["ACC-4471-0092", "ACC-4471-0177"],
      "fields": "standard", // excludes tax identifiers by default
      "expires_at": "2026-11-02T00:00:00Z"
    },
    {
      "type": "brokerage_trade",
      "accounts": ["ACC-4471-0092"],
      "actions": ["propose", "place"],
      "instruments": { "asset_class": ["equity", "etf"],
        "exclude": ["options", "futures", "crypto"] },
      "order_types": ["limit"], // no market orders from an agent
      "notional_cap": { "amount": "10000.00", "currency": "USD",
        "period": "P1D" },
      "order_count_cap": { "count": 5, "period": "P1D" },
      "confirmation": "out_of_band", // non-negotiable; server-enforced
      "expires_at": "2026-09-02T00:00:00Z"
    }
  ]
}
```

What each element buys

ELEMENT	RISK IT REMOVES
accounts	An agent granted access for one account cannot reach another, even though the customer is authenticated for both.
instruments.exclude	Removes whole product classes from the agentic channel. Options and crypto need never be reachable by a planner.
order_types	Limit-only removes the worst outcome of a mispriced or hallucinated order — an unbounded market fill.
notional_cap and order_count_cap	Bounds total loss from any single compromised session to a figure the customer chose, decremented atomically with each order.
confirmation	Carried in the grant so it cannot be omitted by a client, negotiated away by a request, or forgotten by a downstream service.

Two separate values <code>expires_at</code>	Read access can be long-lived because it is low-risk; trade authority expires in weeks. Consent is not one undifferentiated grant.
---	---

Note that none of these constraints can be expressed in an OAuth scope string, and none of them can be enforced by a component that only receives a tool name and a set of arguments.

APPENDIX C

Glossary

Agent Skills	Open standard released December 2025 for packaging instructions, templates and scripts as a <code>SKILL.md</code> directory that agents load by progressive disclosure. Adopted across roughly forty clients.
CIBA	Client-Initiated Backchannel Authentication. Decoupled authorisation: the request originates in one channel and is approved in another, typically a push notification to a trusted device.
DPoP (RFC 9449)	Demonstration of Proof-of-Possession. Binds an access token to a client-held key so that a stolen token is unusable without the key.
FAPI 2.0	OpenID Foundation Financial-grade API Security Profile. Mandates sender-constrained tokens and PAR; formally analysed; the appropriate baseline for money-moving APIs.
ID-JAG / XAA	Identity Assertion JWT Authorization Grant, also called Cross-App Access. Lets an enterprise identity provider issue scoped, short-lived tokens for another application's API. Became a stable MCP extension in June 2026; requires a corporate IdP.
Intent receipt	Our term for the signed record binding customer, agent, grant, parameters, confirmation and outcome for a single acting call.
MCP	Model Context Protocol. Introduced November 2024; standardises how AI applications connect to external tools and data. Current authorisation specification 2025-11-25.
PAR (RFC 9126)	Pushed Authorization Requests. Moves the authorisation request into a back-channel call, enabling rich structured grant requests.
PDP / PEP	Policy Decision Point and Policy Enforcement Point. The PDP evaluates policy; the PEP — here, the gateway — enforces the verdict and its obligations.
RAR (RFC 9396)	Rich Authorization Requests. Replaces the flat scope string with a structured <code>authorization_details</code> array capable of expressing a transaction.
Resource Indicators (RFC 8707)	Lets a client name the resource server a token is for, so that resource servers can validate audience and refuse tokens minted for someone else.
Token exchange (RFC 8693)	Exchanges one token for another, typically to cross a trust boundary while preserving the acting party in an <code>act</code> claim.
Tool poisoning	Hiding instructions in a tool's description, schema or return value so that a model reading them changes its behaviour.
WORM	Write Once Read Many. Storage mode required for regulated books-and-records retention.

APPENDIX D

References

Primary sources consulted, current as at August 2026. Where a claim in this report rests on a single source, that source is named in the text.

PROTOCOL AND SPECIFICATION

1. Model Context Protocol — [Security Best Practices](#), specification 2025-11-25. modelcontextprotocol.io/specification/2025-11-25/basic/security_best_practices
2. Model Context Protocol — [Authorization](#), specifications 2025-06-18 and 2025-11-25. modelcontextprotocol.io/specification/2025-06-18/basic/authorization
3. Model Context Protocol — [Enterprise-Managed Authorization extension](#), stable June 2026. modelcontextprotocol.io/extensions/auth/enterprise-managed-authorization
4. IETF — [Identity Assertion JWT Authorization Grant](#), draft-ietf-oauth-identity-assertion-authz-grant.

SECURITY GUIDANCE

1. OWASP — [MCP Security Cheat Sheet](#), Cheat Sheet Series. cheatsheetseries.owasp.org/cheatsheets/MCP_Security_Cheat_Sheet.html
2. OWASP GenAI Security Project — [Top 10 for Agentic Applications 2026](#), published December 2025. [ASI01 Agent Goal Hijack through ASI10 Rogue Agents](#). genai.owasp.org
3. OWASP — [AI Agent Security Cheat Sheet and LLM Prompt Injection Prevention Cheat Sheet](#).
4. Snyk / Koi Security — [Malicious MCP server on npm: postmark-mcp harvests emails](#), September 2025.

STANDARDS

1. OpenID Foundation — [FAPI 2.0 Security Profile](#), final. openid.net/specs/fapi-security-profile-2_0-final.html
2. RFC 9396 — [OAuth 2.0 Rich Authorization Requests](#), May 2023.
3. RFC 8707 — [Resource Indicators for OAuth 2.0](#). RFC 9728 — [OAuth 2.0 Protected Resource Metadata](#).
4. RFC 9449 — [OAuth 2.0 Demonstrating Proof of Possession \(DPoP\)](#). RFC 9126 — [Pushed Authorization Requests](#).
5. RFC 8693 — [OAuth 2.0 Token Exchange](#). RFC 9700 — [OAuth 2.0 Security Best Current Practice](#).

REGULATORY

1. FINRA — [2026 Annual Regulatory Oversight Report](#), GenAI: Continuing and Emerging Trends, December 2025. First standalone GenAI section; includes *Emerging Trends in GenAI: Agents*. finra.org/rules-guidance/guidance/reports/2026-finra-annual-regulatory-oversight-report/gen-ai
2. FINRA — [Regulatory Notice 24-09](#): obligations when using generative artificial intelligence and large language models.
3. FINRA — [Regulatory Notice 21-29](#): supervisory obligations related to outsourcing to third-party vendors.
4. FINRA Rules 3110 (Supervision) and 4511 (General Requirements, books and records); SEA Rule 17a-4.
5. SEC — [Regulation S-P: Privacy of Consumer Financial Information and Safeguarding Customer Information](#), Release 34-100155, adopted 16 May 2024; compliance December 2025 / June 2026.
6. NIST — [Artificial Intelligence Risk Management Framework \(AI RMF 1.0\)](#), January 2023.

AGENT CAPABILITY PACKAGING

1. Anthropic — [Agent Skills open standard and SKILL .md format](#), released December 2025.